

Module 1 - Data Ingestion And Exploration

April 19, 2025

```
[1]: from pyspark.sql import SparkSession
      spark = SparkSession.builder \
        .appName('OlistData') \
        .getOrCreate()
```

25/04/12 15:54:45 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.

```
[2]: spark
```

```
[2]: <pyspark.sql.session.SparkSession at 0x7f7610b34cd0>
```

```
[3]: !hadoop fs -ls /data/olist/
```

Found 9 items

```
-rw-r--r--  2 klpd1299 hadoop    9033957 2025-04-12 15:51
/data/olist/olist_customers_dataset.csv
-rw-r--r--  2 klpd1299 hadoop    61273883 2025-04-12 15:51
/data/olist/olist_geolocation_dataset.csv
-rw-r--r--  2 klpd1299 hadoop    15438671 2025-04-12 15:51
/data/olist/olist_order_items_dataset.csv
-rw-r--r--  2 klpd1299 hadoop     5777138 2025-04-12 15:51
/data/olist/olist_order_payments_dataset.csv
-rw-r--r--  2 klpd1299 hadoop    14451670 2025-04-12 15:51
/data/olist/olist_order_reviews_dataset.csv
-rw-r--r--  2 klpd1299 hadoop    17654914 2025-04-12 15:51
/data/olist/olist_orders_dataset.csv
-rw-r--r--  2 klpd1299 hadoop     2379446 2025-04-12 15:51
/data/olist/olist_products_dataset.csv
-rw-r--r--  2 klpd1299 hadoop      174703 2025-04-12 15:51
/data/olist/olist_sellers_dataset.csv
-rw-r--r--  2 klpd1299 hadoop         2613 2025-04-12 15:51
/data/olist/product_category_name_translation.csv
```

```
[4]: hdfs_path = '/data/olist/'
```

```
[6]: customer_df = spark.read.csv(hdfs_path + 'olist_customers_dataset.
    ↪ csv', header=True, inferSchema=True)
```

```
[7]: customer_df.show(5)
```

```
[Stage 2:> (0 + 1) / 1]
+-----+-----+-----+-----+
| customer_id | customer_unique_id | customer_zip_code_prefix |
customer_city | customer_state |
+-----+-----+-----+-----+
| 06b8999e2fba1a1fb... | 861eff4711a542e4b... | 14409 |
franca | SP |
| 18955e83d337fd6b2... | 290c77bc529b7ac93... | 9790 |
do c... | SP |
| 4e7b3e00288586ebd... | 060e732b5b29e8181... | 1151 |
sao paulo | SP |
| b2b6027bc5c5109e5... | 259dac757896d24d7... | 8775 |
das cruzeiras | SP |
| 4f2d8ab171c80ec83... | 345ecd01c38d18a90... | 13056 |
campinas | SP |
+-----+-----+-----+-----+
only showing top 5 rows
```

```
[8]: category_translation_df = spark.read.csv(hdfs_path +
↳ 'product_category_name_translation.csv', header=True, inferSchema=True)
orders_df = spark.read.csv(hdfs_path + 'olist_orders_dataset.
↳ csv', header=True, inferSchema=True)
order_item_df = spark.read.csv(hdfs_path + 'olist_order_items_dataset.
↳ csv', header=True, inferSchema=True)
payments_df = spark.read.csv(hdfs_path + 'olist_order_payments_dataset.
↳ csv', header=True, inferSchema=True)
reviews_df = spark.read.csv(hdfs_path + 'olist_order_reviews_dataset.
↳ csv', header=True, inferSchema=True)
products_df = spark.read.csv(hdfs_path + 'olist_products_dataset.
↳ csv', header=True, inferSchema=True)
sellers_df = spark.read.csv(hdfs_path + 'olist_sellers_dataset.
↳ csv', header=True, inferSchema=True)
geolocation_df = spark.read.csv(hdfs_path + 'olist_geolocation_dataset.
↳ csv', header=True, inferSchema=True)
```

```
[10]: customer_df.printSchema()
```

```
root
```

```

|-- customer_id: string (nullable = true)
|-- customer_unique_id: string (nullable = true)
|-- customer_zip_code_prefix: integer (nullable = true)
|-- customer_city: string (nullable = true)
|-- customer_state: string (nullable = true)

```

```
[11]: orders_df.printSchema()
```

```

root
 |-- order_id: string (nullable = true)
 |-- customer_id: string (nullable = true)
 |-- order_status: string (nullable = true)
 |-- order_purchase_timestamp: timestamp (nullable = true)
 |-- order_approved_at: timestamp (nullable = true)
 |-- order_delivered_carrier_date: timestamp (nullable = true)
 |-- order_delivered_customer_date: timestamp (nullable = true)
 |-- order_estimated_delivery_date: timestamp (nullable = true)

```

```
[12]: # Data leakage or drop
print(f'Customers : {customer_df.count()} rows')
print(f'Orders : {orders_df.count()} rows')
```

```

Customers : 99441 rows
Orders : 99441 rows

```

```
[15]: # NULL or say duplicate values

from pyspark.sql.functions import col

customer_df.select([col(c).isNull().alias(c) for c in customer_df.columns]).
    <=>show()
```

```

+-----+-----+-----+-----+-----+
-----+
|customer_id|customer_unique_id|customer_zip_code_prefix|customer_city|customer_
state|
+-----+-----+-----+-----+-----+
-----+
|      false|           false|           false|           false|           false|
false|
|      false|           false|           false|           false|           false|
false|
|      false|           false|           false|           false|           false|
false|
|      false|           false|           false|           false|           false|
false|
|      false|           false|           false|           false|           false|
false|

```

```

false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
|      false|                false|                false|                false|
false|
+-----+-----+-----+-----+-----+
-----+
only showing top 20 rows

```

[16]: *# NULL or say duplicate values*

```

from pyspark.sql.functions import col,when,count

```

```

customer_df.select([count(when(col(c).isNull(),1)).alias(c) for c in_
↳customer_df.columns])).show()

```

```

+-----+-----+-----+-----+-----+
-----+
|customer_id|customer_unique_id|customer_zip_code_prefix|customer_city|customer_
state|
+-----+-----+-----+-----+-----+

```

```

-----+
|           0|           0|           0|           0|
0|
+-----+-----+-----+-----+
-----+

```

```
[17]: # Duplicate values
customer_df.groupBy('customer_id').count().filter('count>1').show()
```

```
[Stage 29:=====> (1 + 1) / 2]
```

```

+-----+-----+
|customer_id|count|
+-----+-----+
+-----+-----+

```

```
[19]: ## customer distribute by state
customer_df.groupBy('customer_state').count().orderBy('count',ascending=True).
↪show()
```

```

+-----+-----+
|customer_state|count|
+-----+-----+
|           RR|   46|
|           AP|   68|
|           AC|   81|
|           AM|  148|
|           RO|  253|
|           TO|  280|
|           SE|  350|
|           AL|  413|
|           RN|  485|
|           PI|  495|
|           PB|  536|
|           MS|  715|
|           MA|  747|
|           MT|  907|
|           PA|  975|
|           CE| 1336|
|           PE| 1652|
|           GO| 2020|
|           ES| 2033|
|           DF| 2140|
+-----+-----+

```

only showing top 20 rows

```
[21]: orders_df.groupBy('order_status').count().orderBy('count',ascending=True).show()
```

```
+-----+-----+
|order_status|count|
+-----+-----+
|    approved|    2|
|    created|    5|
| processing|  301|
|   invoiced|  314|
| unavailable|  609|
|   canceled|  625|
|    shipped| 1107|
|   delivered|96478|
+-----+-----+
```

```
[22]: payments_df.groupBy('payment_type').count().orderBy('count',ascending=True).
      ↪show()
```

```
+-----+-----+
|payment_type|count|
+-----+-----+
| not_defined|    3|
| debit_card| 1529|
|    voucher| 5775|
|    boleto|19784|
| credit_card|76795|
+-----+-----+
```

```
[23]: from pyspark.sql.functions import sum

top_products = order_item_df.groupBy("product_id").agg(sum('price').
      ↪alias('total_sales'))
top_products.orderBy('total_sales',ascending=False).show(20)
```

[Stage 46:=====> (1 + 1) / 2]

```
+-----+-----+
|      product_id|      total_sales|
+-----+-----+
|bb50f2e236e5eea01...|      63885.0|
|6cdd53843498f9289...| 54730.200000000005|
|d6160fb7873f18409...|48899.340000000004|
|d1c427060a0f73f6b...| 47214.510000000006|
|99a4788cb24856965...|43025.5600000000085|
|3dd2a17168ec895c7...| 41082.600000000005|
|25c38557cf793876c...| 38907.320000000001|
|5f504b3a1c75b73d6...|37733.899999999994|
```

```
|53b36df67ebb7c415...| 37683.420000000001|
|aca2eb7d00ea1a7b8...| 37608.900000000007|
|e0d64dcfaa3b6db5c...|          31786.82|
|d285360f29ac7fd97...|31623.809999999983|
|7a10781637204d8d1...|          30467.5|
|f1c7f353075ce59d8...|          29997.36|
|f819f0c84a64f02d3...|29024.479999999996|
|588531f8ec37e7d5f...|28291.989999999998|
|422879e10f4668299...|26577.219999999972|
|16c4e87b98a9370a9...|          25034.0|
|5a848e4ab52fd5445...|24229.029999999962|
|a62e25e09e05e6faf...|          24051.0|
+-----+-----+
only showing top 20 rows
```

```
[31]: #Average Delivery Time Analysis
delivery_df = orders_df.
        ↪select('order_id','order_purchase_timestamp','order_delivered_customer_date').
        ↪show()
```

```
+-----+-----+-----+
|          order_id|order_purchase_timestamp|order_delivered_customer_date|
+-----+-----+-----+
|e481f51cbdc54678b...|      2017-10-02 10:56:33|      2017-10-10 21:25:13|
|53cdb2fc8bc7dce0b...|      2018-07-24 20:41:37|      2018-08-07 15:27:45|
|47770eb9100c2d0c4...|      2018-08-08 08:38:49|      2018-08-17 18:06:29|
|949d5b44dbf5de918...|      2017-11-18 19:28:06|      2017-12-02 00:28:42|
|ad21c59c0840e6cb8...|      2018-02-13 21:18:39|      2018-02-16 18:17:02|
|a4591c265e18cb1dc...|      2017-07-09 21:57:05|      2017-07-26 10:57:55|
|136cce7faa42fdb2c...|      2017-04-11 12:22:08|              NULL|
|6514b8ad8028c9f2c...|      2017-05-16 13:10:30|      2017-05-26 12:55:51|
|76c6e866289321a7c...|      2017-01-23 18:29:09|      2017-02-02 14:08:10|
|e69bfb5eb88e0ed6a...|      2017-07-29 11:55:02|      2017-08-16 17:14:30|
|e6ce16cb79ec1d90b...|      2017-05-16 19:41:10|      2017-05-29 11:18:31|
|34513ce0c4fab462a...|      2017-07-13 19:58:11|      2017-07-19 14:04:48|
|82566a660a982b15f...|      2018-06-07 10:06:19|      2018-06-19 12:05:52|
|5ff96c15d0b717ac6...|      2018-07-25 17:44:10|      2018-07-30 15:52:25|
|432aaf21d85167c2c...|      2018-03-01 14:14:28|      2018-03-12 23:36:26|
|dcb36b511fcac050b...|      2018-06-07 19:03:12|      2018-06-21 15:34:32|
|403b97836b0c04a62...|      2018-01-02 19:00:43|      2018-01-20 01:38:59|
|116f0b09343b49556...|      2017-12-26 23:41:31|      2018-01-08 22:36:36|
|85ce859fd6dc634de...|      2017-11-21 00:03:41|      2017-11-27 18:28:00|
|83018ec114eee8641...|      2017-10-26 15:54:26|      2017-11-08 22:22:00|
+-----+-----+-----+
only showing top 20 rows
```

[]: